

Database Design

Team members: Sally Xue, Pradyumann Singhal, Lily Zhang, Joseph Schanne

GCP Connection Screenshot

for implementing the database tables locally or on GCP, you should provide a screenshot of the connection (i.e. showing your terminal/command-line information)

```
mysql> xuesally3@cloudshell:~ (cs411-team109)$ gcloud sql connect cs411-final-project --user=root
Allowlisting your IP for incoming connection for 5 minutes...done.
Connecting to database with SQL user [root].Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 7328
Server version: 8.0.31-google (Google)

Copyright (c) 2000, 2024, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> use spotify_comment_hub
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> SHOW TABLES;
+-----+
| Tables_in_spotify_comment_hub |
+-----+
| ALBUMS                         |
| ARTISTS                       |
| COMMENTS                      |
| SONGS                         |
| USERS                         |
+-----+
5 rows in set (0.00 sec)
```

DDL Commands

```
CREATE TABLE ARTISTS (
    ArtistID VARCHAR(255) PRIMARY KEY,
    ArtistName VARCHAR(255)
)
```

```

CREATE TABLE USERS (
    UserID VARCHAR(255) PRIMARY KEY,
    Username VARCHAR(255),
    Password VARCHAR(255),
    Email VARCHAR(255)
);

CREATE TABLE ALBUMS (
    AlbumID VARCHAR(255) PRIMARY KEY,
    ArtistID VARCHAR(255),
    AlbumName VARCHAR(255),
    Description VARCHAR(255),
    FOREIGN KEY (ArtistID) REFERENCES ARTISTS(ArtistID)
);

CREATE TABLE SONGS (
    SongID VARCHAR(255) PRIMARY KEY,
    SongName VARCHAR(255),
    ArtistID VARCHAR(255),
    AlbumID VARCHAR(255),
    ReleaseDate DATE,
    FOREIGN KEY (ArtistID) REFERENCES ARTISTS(ArtistID),
    FOREIGN KEY (AlbumID) REFERENCES ALBUMS(AlbumID)
);

CREATE TABLE COMMENTS (
    CommentID VARCHAR(255) PRIMARY KEY,
    UserID VARCHAR(255),
    SongID VARCHAR(255),
    CommentInfo VARCHAR(255),
    Rating INTEGER,
    CreatedOn DATE,
    ResponseTo VARCHAR(255),
    FOREIGN KEY (UserID) REFERENCES USERS(UserID),
    FOREIGN KEY (SongID) REFERENCES SONGS(SongID)
);

```

Table Row Count

ALBUMS:

```
mysql> SELECT COUNT(*) FROM ALBUMS;
+-----+
| COUNT(*) |
+-----+
|      1631 |
+-----+
1 row in set (0.01 sec)
```

ARTISTS:

```
mysql> SELECT COUNT(*) FROM ARTISTS;
+-----+
| COUNT(*) |
+-----+
|      1104 |
+-----+
1 row in set (0.00 sec)
```

SONGS:

```
mysql> SELECT COUNT(*) FROM SONGS;
+-----+
| COUNT(*) |
+-----+
|      1780 |
+-----+
1 row in set (0.00 sec)
```

Advanced Queries with Screenshots

- **Song Count per Artist**

SELECT COUNT(SongName), a.ArtistName FROM SONGS s NATURAL JOIN ARTISTS a GROUP BY ArtistID LIMIT 15;

```
mysql> SELECT COUNT(SongName), a.ArtistName FROM SONGS s NATURAL JOIN ARTISTS a GROUP BY ArtistID LIMIT 15;
+-----+-----+
| COUNT(SongName) | ArtistName |
+-----+-----+
| 1 | Blue Öyster Cult |
| 1 | Lupe Fiasco |
| 1 | Adriatique |
| 1 | Zoe Wees |
| 1 | Lavern |
| 2 | CKay |
| 5 | Maroon 5 |
| 1 | El Padrinito Toys |
| 1 | bby |
| 1 | Brandy |
| 1 | BLK ODYSSEY |
| 9 | Taylor Swift |
| 3 | KISS |
| 2 | Jason Derulo |
| 1 | Joaquina |
+-----+-----+
15 rows in set (0.00 sec)
```

- **Album Count per Artist**

SELECT COUNT(AlbumName), a.ArtistName FROM ALBUMS ab NATURAL JOIN ARTISTS a GROUP BY ArtistID;

```
mysql> SELECT COUNT(AlbumName), a.ArtistName FROM ALBUMS ab NATURAL JOIN ARTISTS a GROUP BY ArtistID LIMIT 15;
+-----+-----+
| COUNT(AlbumName) | ArtistName |
+-----+-----+
| 1 | Blue Öyster Cult |
| 1 | Lupe Fiasco |
| 1 | Adriatique |
| 1 | Zoe Wees |
| 1 | Lavern |
| 2 | CKay |
| 3 | Maroon 5 |
| 1 | El Padrinito Toys |
| 1 | bby |
| 1 | Brandy |
| 1 | BLK ODYSSEY |
| 6 | Taylor Swift |
| 3 | KISS |
| 2 | Jason Derulo |
| 1 | Joaquina |
+-----+-----+
15 rows in set (0.01 sec)
```

- **Give User ID, ArtistID, Count of Comments**

SELECT UserID, ArtistID, COUNT(CommentID) as numComments FROM COMMENTS NATURAL JOIN SONGS NATURAL JOIN ARTISTS

GROUP BY UserID, ArtistID ORDER BY numComments DESC LIMIT 15;

```
mysql> SELECT UserID, ArtistID, COUNT(CommentID) as numComments FROM  
-> COMMENTS NATURAL JOIN SONGS NATURAL JOIN ARTISTS  
-> GROUP BY UserID, ArtistID ORDER BY numComments DESC LIMIT 15;
```

UserID	ArtistID	numComments
0	0tbeZu9lv8YEKSQ9tZSslu	3
1	3VVLqeEqQQqTgT8YhfY9Z6	2
0	45lcbTsX07JWzmTIjcdyBz	2
1	45lcbTsX07JWzmTIjcdyBz	2
0	Op4nmQO2msCgU4IF37Wi3j	2
1	Op4nmQO2msCgU4IF37Wi3j	2
0	3tJoFztHeIJkJWMrx0td2f	2
1	3tJoFztHeIJkJWMrx0td2f	2
0	3VVLqeEqQQqTgT8YhfY9Z6	2
1	0tbeZu9lv8YEKSQ9tZSslu	1
0	3hv9jJF3adDNsBSIQDqcjp	1
1	3hv9jJF3adDNsBSIQDqcjp	1
0	1Xv1qZHJ1hnRlWHRTZ3uci	1
1	1Xv1qZHJ1hnRlWHRTZ3uci	1
0	4Ge9GwmWnOQsohwPTrXyHc	1

15 rows in set (0.01 sec)

- **Get Avg Ratings of all artists based on their songs**

SELECT ArtistID, Avg(Rating) as avgRating, COUNT(Rating) as
NumRatings FROM COMMENTS NATURAL JOIN SONGS NATURAL
JOIN ARTISTS GROUP BY ArtistID LIMIT 15;

```
mysql> SELECT ArtistID, Avg(Rating) as avgRating, COUNT(Rating) as NumRatings FROM COMMENTS NATURAL JOIN SONGS NATURAL JOIN ARTISTS GROUP BY ArtistID LIMIT 15;
```

ArtistID	avgRating	NumRatings
0tbeZu9lv8YEKSQ9tZSslu	8.5000	4
3VVLqeEqQQqTgT8YhfY9Z6	6.2500	4
45lcbTsX07JWzmTIjcdyBz	6.7500	4
Op4nmQO2msCgU4IF37Wi3j	7.7500	4
3tJoFztHeIJkJWMrx0td2f	7.0000	4
3hv9jJF3adDNsBSIQDqcjp	5.5000	2
1Xv1qZHJ1hnRlWHRTZ3uci	5.0000	2
4Ge9GwmWnOQsohwPTrXyHc	5.0000	2
24DO0PijjITGIEWs08XaPs	5.0000	2
2SmW1lF1Bjn4IfBzB2D1Sh	6.0000	2
6EP1BSH2RSiettcz1z7ihV	4.5000	2
0cmWgDlu9CwTgxPhf403hb	4.5000	2
2qoQgPaileR0KCwE2Y8wOG	5.0000	2
4q3ewBCX7sLwd24euuV69X	5.0000	2
5ObBtv5VunwwhQaXXnUrsM	5.0000	2

15 rows in set (0.01 sec)

Indexing Analysis

- **Song Count per Artist**

```
SELECT COUNT(SongName), a.ArtistName, s.AlbumID FROM SONGS s  
NATURAL JOIN ARTISTS a GROUP BY ArtistID;
```

Default index:

```
| -> Table scan on <temporary> (actual time=9.769..9.898 rows=1099 loops=1)  
-> Aggregate using temporary table (actual time=9.767..9.767 rows=1099 loops=1)  
-> Nested loop inner join (cost=737.48 rows=1788) (actual time=1.745..8.336 rows=1780 loops=1)  
-> Table scan on a (cost=111.65 rows=1104) (actual time=0.602..0.962 rows=1104 loops=1)  
-> Index lookup on s using ArtistID (ArtistID=a.ArtistID) (cost=0.41 rows=2) (actual time=0.005..0.005 rows=2 loops=1104)  
|
```

Analysis: The costliest part of this query is the nested loop join (row 3), which is the NATURAL JOIN part in the SQL query. Specifically the index lookup on s using ArtistID, where each row from a requires an index lookup on s, increasing the join cost to 737.48.

Indexing 1: CREATE INDEX idx_songs_artist ON SONGS (ArtistID);

```
mysql> EXPLAIN ANALYZE SELECT COUNT(SongName), a.ArtistName FROM SONGS s NATURAL JOIN ARTISTS a GROUP BY ArtistID;  
+-----+  
| EXPLAIN |  
+-----+  
|  
+-----+  
| -> Table scan on <temporary> (actual time=6.097..6.226 rows=1099 loops=1)  
-> Aggregate using temporary table (actual time=6.095..6.095 rows=1099 loops=1)  
-> Nested loop inner join (cost=738.15 rows=1790) (actual time=0.081..4.877 rows=1780 loops=1)  
-> Table scan on a (cost=111.75 rows=1105) (actual time=0.050..0.297 rows=1104 loops=1)  
-> Index lookup on s using fk_artist_id (ArtistID=a.ArtistID) (cost=0.41 rows=2) (actual time=0.003..0.004 rows=2 loops=1104)  
|  
+-----+  
1 row in set (0.01 sec)
```

The index `idx_songs_artist` on the `ArtistID` column in the `SONGS` table does not optimize the cost in this query due to the nature of the NATURAL JOIN operation. While indexing `ArtistID` can improve the lookup speed for rows in the `SONGS` table when matched with the `ARTISTS` table, it still performs numerous index lookups for each row in the `ARTISTS` table, which does not optimize the cost.

Indexing 2: CREATE INDEX idx_artists_name ON ARTISTS(ArtistName);

```
mysql> EXPLAIN ANALYZE SELECT COUNT(SongName), a.ArtistName FROM SONGS s NATURAL JOIN ARTISTS a GROUP BY ArtistID;
+-----+
| EXPLAIN |
+-----+
|         |
+-----+
|         |
+-----+
| -> Table scan on <temporary> (actual time=6.130..6.249 rows=1099 loops=1) |
| -> Aggregate using temporary table (actual time=6.128..6.128 rows=1099 loops=1) |
| -> Nested loop inner join (cost=738.15 rows=1790) (actual time=0.050..4.896 rows=1780 loops=1) |
|   -> Covering index scan on a using idx_artists_name (cost=111.75 rows=1105) (actual time=0.030..0.246 rows=1104 loops=1) |
|   -> Index lookup on s using fk_artist_id (ArtistID=a.ArtistID) (cost=0.41 rows=2) (actual time=0.003..0.004 rows=2 loops=1104) |
+-----+
| row in set (0.01 sec) |
+-----+
```

The index `idx_artists_name` does not optimize the query because it is not utilized in the join or grouping operations. In this query, the join condition is based solely on the `ArtistID`, which is not affected by the `ArtistName` index. Consequently, the database engine will still need to perform a full scan of the `ARTISTS` table to find matching rows for the join, leading to the cost not changing.

Indexing 3: `CREATE INDEX idx_songs_songname ON SONGS (SongName);`

```
mysql> EXPLAIN ANALYZE SELECT COUNT(SongName), a.ArtistName FROM SONGS s NATURAL JOIN ARTISTS a GROUP BY ArtistID;
+-----+
| EXPLAIN |
+-----+
|         |
+-----+
|         |
+-----+
| -> Table scan on <temporary> (actual time=6.097..6.226 rows=1099 loops=1) |
| -> Aggregate using temporary table (actual time=6.095..6.095 rows=1099 loops=1) |
| -> Nested loop inner join (cost=738.15 rows=1790) (actual time=0.081..4.877 rows=1780 loops=1) |
|   -> Table scan on a (cost=111.75 rows=1105) (actual time=0.050..0.297 rows=1104 loops=1) |
|   -> Index lookup on s using fk_artist_id (ArtistID=a.ArtistID) (cost=0.41 rows=2) (actual time=0.003..0.004 rows=2 loops=1104) |
+-----+
| row in set (0.01 sec) |
+-----+
```

Similar to indexing 2, the `SongName` column in the `SONGS` table does not optimize the query because it is not involved in the join or the aggregation process of the query.

- **Album Count per Artist**

```
SELECT COUNT(AlbumName), a.ArtistName
FROM ALBUMS ab NATURAL JOIN ARTISTS a
GROUP BY ArtistID;
```

Default index:

```
mysql> explain analyze SELECT COUNT(AlbumName), a.ArtistName FROM ALBUMS ab NATURAL JOIN ARTISTS a GROUP BY ArtistID;
+-----+
| EXPLAIN
+-----+

-> Table scan on <temporary> (actual time=6.127..6.244 rows=1099 loops=1)
-> Aggregate using temporary table (actual time=6.125..6.125 rows=1099 loops=1)
-> Nested loop inner join (cost=685.72 rows=1640) (actual time=0.077..4.899 rows=1631 loops=1)
-> Table scan on a (cost=111.75 rows=1105) (actual time=0.047..0.307 rows=1104 loops=1)
-> Index lookup on ab using ArtistID (ArtistID=a.ArtistID) (cost=0.37 rows=1) (actual time=0.003..0.004 rows=1 loops=1104)
```

Indexing 1: CREATE INDEX idx_albums_artist ON ALBUMS (ArtistID);

```
mysql> CREATE INDEX idx_albums_artist ON ALBUMS (ArtistID);
Query OK, 0 rows affected (0.09 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> EXPLAIN ANALYZE SELECT COUNT(AlbumName), a.ArtistName
-> FROM ALBUMS ab NATURAL JOIN ARTISTS a
-> GROUP BY ArtistID;
+-----+
| EXPLAIN
+-----+

-> Table scan on <temporary> (actual time=6.348..6.476 rows=1099 loops=1)
-> Aggregate using temporary table (actual time=6.347..6.347 rows=1099 loops=1)
-> Nested loop inner join (cost=685.72 rows=1640) (actual time=0.583..5.270 rows=1631 loops=1)
-> Table scan on a (cost=111.75 rows=1105) (actual time=0.543..0.781 rows=1104 loops=1)
-> Index lookup on ab using idx_albums_artist (ArtistID=a.ArtistID) (cost=0.37 rows=1) (actual time=0.003..0.004 rows=1 loops=1104)

1 row in set (0.01 sec)
```

The index `idx_albums_artist` on the `ArtistID` column in the `ALBUMS` table does not optimize the cost in this query due to the nature of the `NATURAL JOIN` operation. While indexing `ArtistID` can improve the lookup speed for rows in the `ALBUMS` table when matched with the `ARTISTS` table, it still performs numerous index lookups for each row in the `ARTISTS` table, which does not optimize the cost.

Indexing 2: CREATE INDEX idx_albums_name ON ALBUMS (AlbumName);

```
-> Table scan on <temporary> (actual time=5.646..5.763 rows=1099 loops=1)
-> Aggregate using temporary table (actual time=5.644..5.644 rows=1099 loops=1)
-> Nested loop inner join (cost=685.72 rows=1640) (actual time=0.087..4.582 rows=1631 loops=1)
-> Table scan on a (cost=111.75 rows=1105) (actual time=0.056..0.294 rows=1104 loops=1)
-> Index lookup on ab using ArtistID (ArtistID=a.ArtistID) (cost=0.37 rows=1) (actual time=0.003..0.004 rows=1 loops=1104)

1 row in set (0.01 sec)
```

The index `idx_albums_name` does not optimize the query because it is not utilized in the join or grouping operations. In this query, the join condition is based solely on the `ArtistID`, which is not affected by the `AlbumName` index. Consequently, the database engine will still need to perform a full scan of the `ARTISTS` table to find matching rows for the join, leading to the cost not changing.

Indexing 3: CREATE INDEX idx_artists_artistname ON ARTISTS (ArtistName);

```
-----+
-> Table scan on <temporary> (actual time=5.594..5.710 rows=1099 loops=1)
-> Aggregate using temporary table (actual time=5.592..5.592 rows=1099 loops=1)
-> Nested loop inner join (cost=685.72 rows=1640) (actual time=0.067..4.559 rows=1631 loops=1)
-> Covering index scan on a using idx_artists_artistname (cost=111.75 rows=1105) (actual time=0.029..0.227 rows=1104 loops=1)
-> Index lookup on ab using ArtistID (ArtistID=a.ArtistID) (cost=0.37 rows=1) (actual time=0.003..0.004 rows=1 loops=1104)
|
-----+

```

Similar to indexing 2, the **ArtistName** column in the **ARTISTS** table does not optimize the query because it is not involved in the join or the aggregation process of the query.

- **Give User ID, ArtistID, Count of Comments**

```
SELECT UserID, ArtistID, COUNT(CommentID) as numComments
FROM COMMENTS NATURAL JOIN SONGS NATURAL JOIN ARTISTS
GROUP BY UserID, ArtistID ORDER BY numComments DESC;
```

Default Index:

```
-----+
-> Sort: numComments DESC (actual time=0.339..0.341 rows=38 loops=1)
-> Table scan on <temporary> (actual time=0.279..0.320 rows=38 loops=1)
-> Aggregate using temporary table (actual time=0.278..0.278 rows=38 loops=1)
-> Nested loop inner join (cost=38.65 rows=48) (actual time=0.070..0.222 rows=48 loops=1)
-> Nested loop inner join (cost=21.85 rows=48) (actual time=0.064..0.161 rows=48 loops=1)
-> Filter: (COMMENTS.SongID is not null) (cost=5.05 rows=48) (actual time=0.041..0.056 rows=48 loops=1)
-> Table scan on COMMENTS (cost=5.05 rows=48) (actual time=0.040..0.052 rows=48 loops=1)
-> Filter: (SONGS.ArtistID is not null) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
-> Single-row index lookup on SONGS using PRIMARY (SongID=COMMENTS.SongID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
-> Single-row covering index lookup on ARTISTS using PRIMARY (ArtistID=SONGS.ArtistID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=48)
|
-----+

```

Indexing 1: CREATE INDEX idx_comments_userid ON COMMENTS (UserID);

```
mysql> CREATE INDEX idx_comments_userid ON COMMENTS (UserID);
Query OK, 0 rows affected (0.09 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> EXPLAIN ANALYZE SELECT UserID, ArtistID, COUNT(CommentID) as numComments
-> FROM COMMENTS NATURAL JOIN SONGS NATURAL JOIN ARTISTS
-> GROUP BY UserID, ArtistID ORDER BY numComments DESC;
+-----+
| EXPLAIN |
+-----+
|         |
+-----+
|         |
+-----+
-> Sort: numComments DESC (actual time=0.279..0.281 rows=38 loops=1)
-> Table scan on <temporary> (actual time=0.259..0.262 rows=38 loops=1)
-> Aggregate using temporary table (actual time=0.257..0.257 rows=38 loops=1)
-> Nested loop inner join (cost=38.65 rows=48) (actual time=0.069..0.206 rows=48 loops=1)
-> Nested loop inner join (cost=21.85 rows=48) (actual time=0.062..0.147 rows=48 loops=1)
-> Filter: (COMMENTS.SongID is not null) (cost=5.05 rows=48) (actual time=0.040..0.053 rows=48 loops=1)
-> Table scan on COMMENTS (cost=5.05 rows=48) (actual time=0.039..0.049 rows=48 loops=1)
-> Filter: (SONGS.ArtistID is not null) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
-> Single-row index lookup on SONGS using PRIMARY (SongID=COMMENTS.SongID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
-> Single-row covering index lookup on ARTISTS using PRIMARY (ArtistID=SONGS.ArtistID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=48)
|
+-----+
1 row in set (0.00 sec)

```

This is mainly due to the NATURAL JOIN nature, similar to previous Indexing 1s for other queries.

Indexing 2: CREATE INDEX idx_songs_artistid ON SONGS (ArtistID);


```

-> Table scan on <temporary> (actual time=0.368..0.373 rows=19 loops=1)
-> Aggregate using temporary table (actual time=0.367..0.367 rows=19 loops=1)
    -> Nested loop inner join (cost=38.65 rows=48) (actual time=0.085..0.286 rows=48 loops=1)
        -> Nested loop inner join (cost=21.85 rows=48) (actual time=0.076..0.205 rows=48 loops=1)
            -> Filter: (COMMENTS.SongID is not null) (cost=5.05 rows=48) (actual time=0.047..0.071 rows=48 loops=1)
                -> Table scan on COMMENTS (cost=5.05 rows=48) (actual time=0.046..0.066 rows=48 loops=1)
            -> Filter: (SONGS.ArtistID is not null) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=48)
                -> Single-row index lookup on SONGS using PRIMARY (SongID=COMMENTS.SongID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
        -> Single-row covering index lookup on ARTISTS using PRIMARY (ArtistID=SONGS.ArtistID) (cost=0.25 rows=1) (actual time=0.001..0.002 rows=1 loops=48)

```

Indexing 1: CREATE INDEX idx_songs_artistid ON SONGS (ArtistID);

```
yysql> EXPLAIN ANALYZE SELECT ArtistID, Avg(Rating) as avgRating, COUNT(Rating) as NumRatings FROM COMMENTS NATURAL JOIN SONGS NATURAL JOIN ARTISTS GROUP BY ArtistID;
```

```
-----+-----
```

```
| EXPLAIN
```

```
-----+-----
```

```
|
```

```
+-----+-----
```

```
-> Table scan on <temporary> (actual time=0.250..0.252 rows=19 loops=1)
```

```
-> Aggregate using temporary table (actual time=0.250..0.250 rows=19 loops=1)
```

```
-> Nested loop inner join (cost=38.65 rows=48) (actual time=0.066..0.204 rows=48 loops=1)
```

```
-> Nested loop inner join (cost=21.65 rows=48) (actual time=0.061..0.147 rows=48 loops=1)
```

```
-> Filter: (COMMENTS.SongID is not null) (cost=5.05 rows=48) (actual time=0.039..0.052 rows=48 loops=1)
```

```
-> Table scan on COMMENTS (cost=5.05 rows=48) (actual time=0.038..0.048 rows=48 loops=1)
```

```
-> Filter: (SONGS.ArtistID is not null) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
```

```
-> Single-row index lookup on SONGS using PRIMARY (SongID=COMMENTS.SongID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
```

```
-> Single-row covering index lookup on ARTISTS using PRIMARY (ArtistID=SONGS.ArtistID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=48)
```

```
+
```

```
1 row in set (0.01 sec)
```

Indexing 2: CREATE INDEX idx_comments_rating ON COMMENTS (Rating);

```
mysql> CREATE INDEX idx_comments_rating ON COMMENTS (Rating);
Query OK, 0 rows affected (0.07 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> EXPLAIN ANALYZE SELECT ArtistID, Avg(Rating) as avgRating, COUNT(Rating) as NumRatings FROM COMMENTS NATURAL JOIN SONGS NATURAL JOIN ARTISTS GROUP BY ArtistID;
```

```
-----+-----
```

```
| EXPLAIN
```

```
-----+-----
```

```
|
```

```
-----+-----
```

```
-> Table scan on <temporary> (actual time=0.378..0.381 rows=19 loops=1)
```

```
-> Aggregate using temporary table (actual time=0.378..0.378 rows=19 loops=1)
```

```
    -> Nested loop inner join (cost=38.65 rows=48) (actual time=0.081..0.306 rows=48 loops=1)
```

```
        -> Nested loop inner join (cost=21.85 rows=48) (actual time=0.074..0.224 rows=48 loops=1)
```

```
            -> Filter: (COMMENTS.SongID is not null) (cost=5.05 rows=48) (actual time=0.051..0.073 rows=48 loops=1)
```

```
                -> Table scan on COMMENTS (cost=5.05 rows=48) (actual time=0.049..0.067 rows=48 loops=1)
```

```
                    -> Filter: (SONGS.ArtistID is not null) (cost=0.25 rows=1) (actual time=0.002..0.003 rows=1 loops=48)
```

```
                        -> Single-row index lookup on SONGS using PRIMARY (SongID=COMMENTS.SongID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=48)
```

```
                            -> Single-row covering index lookup on ARTISTS using PRIMARY (ArtistID=SONGS.ArtistID) (cost=0.25 rows=1) (actual time=0.001..0.002 rows=1 loops=48)
```

```
-----+-----
```

```
1 row in set (0.00 sec)
```

Indexing 3: CREATE INDEX idx_comments_songid ON COMMENTS (SongID);

(We use this because although SongID is not explicit in the query, it is used in the natural join of Song and Comments table)

